

FakeVLM single-linear latency comparison: batch 8

