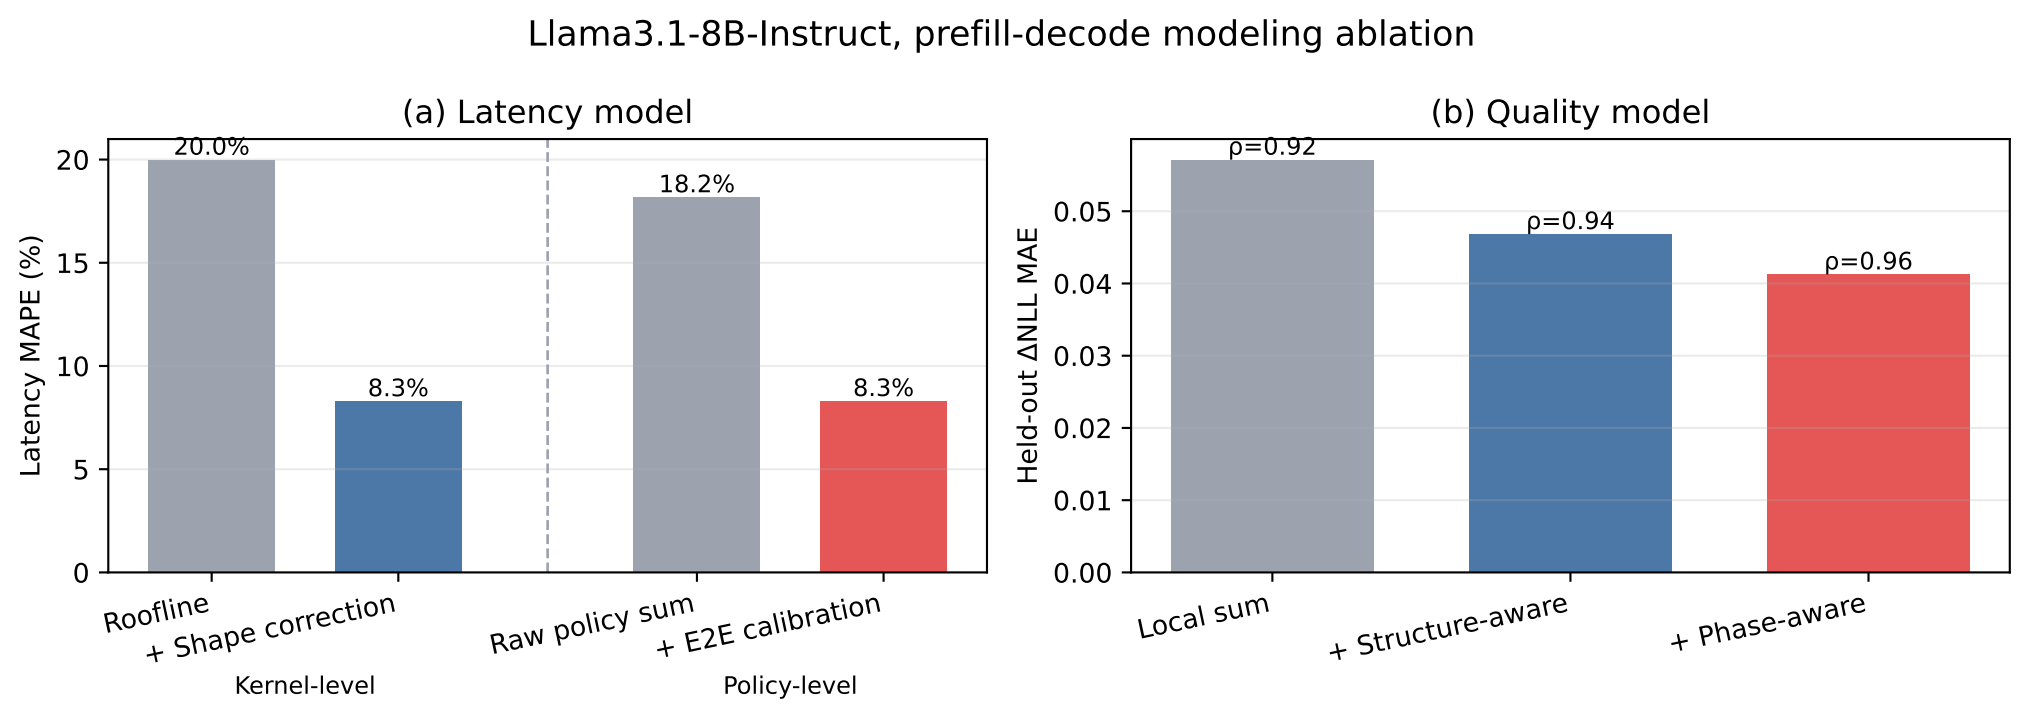
Llama3.1-8B-Instruct, prefill-decode modeling ablation
(a) Latency model
Latency MAPE (%)
20.0%
8.3%
18.2%
8.3%
Roofline
+ Shape correction
Raw policy sum
+ E2E calibration
Kernel-level
Policy-level
(b) Quality model
Held-out ΔNLL MAE
ρ=0.92
ρ=0.94
ρ=0.96
Local sum
+ Structure-aware
+ Phase-aware