

Pure Prefill Latency: batch size=16, input tokens=1024, M=16384 | GPU: NVIDIA RTX 5090 32GB

Prefill Latency (ms)

dense\_bf16 dense\_nvfp4 sparse\_bf16 sparse\_nvfp4 marlin\_nvfp4 hybrid

