



Accuracy vs. Compression Ratio Across Vision Models

Original NVFP4 Rescale Dense baseline (FP32)

NVFP4 + Unstructured Sparse



NVFP4 + 4:8 Structured Sparse



Bar height: ImageNet Top-1 accuracy. CR labels use real file-size ratios from non-Rescale packed checkpoints; Rescale reuses the matching Original ratio.